

2-3 trees

Dictionary: an abstract data type

A container that maps *keys* to *values*

Dictionary operations

- Insert
- Search
- Delete

Several possible implementations

- Balanced search trees
- Hash tables

2-3 trees

A kind of balanced search tree

Assume keys are totally ordered ($<$, $>$, $=$)

Assume n key/value pairs are stored in the dictionary


Time per dictionary operation is $O(\log n)$

Support of other useful operations as well

Basic structure: a tree

- key/value pairs stored only at leaves (no duplicate keys)
- all leaves at the same depth (i.e., distance from root)
- looking at the leaves from left to right, keys appear in sorted order
- each internal node:
 - has either 2 or 3 children
 - has a “guide”: the maximum key in its subtree

Example



Let $h :=$ height of tree (Recall: height = max depth of any node)

Let $n :=$ # of leaves

Claim: $n \geq 2^h$

- Proof by induction on h
- Base case: $h = 0, n = 1$ ✓
- Induction step: $h > 0$, assume claim holds for $h - 1$
 - Tree has a root node, which has either 2 or 3 children
 - Each of these children is the root of a subtree, which itself is a 2-3 tree of height $h - 1$
 - By induction hypothesis, if the i th subtree has n_i leaves, then $n_i \geq 2^{h-1}$ [here, $i = 1 \dots 2$ (or 3)]
 - $\therefore n = \sum_i n_i \geq \sum_i 2^{h-1} \geq 2 \cdot 2^{h-1} = 2^h$ ✓

Corollary: $h \leq \log_2 n$

Example Data Layout (Java syntax)

```
class Node {  
    KeyType guide;  
    // guide points to max key in subtree rooted at node  
}  
  
class InternalNode extends Node {  
    Node child0, child1, child2;  
    // child0 and child1 are always non-null  
    // child2 is null iff node has only 2 children  
}  
  
class LeafNode extends Node {  
    // guide points to the key  
    ValueType value;  
}
```

Search(x): // use guides to search for the key x

$p \leftarrow$ root of tree

$h \leftarrow$ height of tree

repeat h times

 // p points to an internal node

 if $x \leq p.\text{child0}.\text{guide}$ then

$p \leftarrow p.\text{child0}$

 else if $p.\text{child2} = \text{null}$ or $x \leq p.\text{child1}.\text{guide}$ then

$p \leftarrow p.\text{child1}$

 else

$p \leftarrow p.\text{child2}$


 // p now points to a leaf node

 if $x = p.\text{guide}$ then

 return $p.\text{value}$

 else

 return null // or some default value




Search Invariant

Insert(x): Search for x , and if it should belong under p :

add x as a child of p (if not already present)

if p now has 4 children:

- split p into two nodes, p_1 and p_2 , each with two children



- process p 's parent in the same way
- Special case: no parent — create new root, increasing height of tree by 1

Also need to update “guides” — easy

Time = $O(\text{height}) = O(\log n)$


Insert 5




Insert 21




Insert 8





Insert 63




Insert 69




Insert 10




Insert 69







Insert 6



Insert 10



Insert 7




Delete(x): Search for x , and if found under p :


remove x

if p now only has one child:



- if p is the root: delete p (height decreases by 1)
- if one of p 's adjacent siblings has 3 children: p adopts one




- if none of p 's adjacent siblings has 3 children:
 - one sibling q must have 2 children
 - give p 's only child to q
 - delete p
 - process p 's parent




Delete 69




(give)




Delete 8





(give)




(adopt)




(delete root)







Delete 10






Delete 8



(give)



(adopt)



2-3 trees: summary

Assume n keys in dictionary

Running time for lookup, insert, delete:

$O(\log n)$ comparisons, plus $O(\log n)$ overhead

Space: $O(n)$ pointers

Note: in the literature, “traditional” 2-3 trees usually store the guides *in the parent node*

- every node contains one or two guides

“Traditional” 2-3 Trees

Idea: move the guides into the parent node

```
class Node { }  
class InternalNode extends Node {  
    KeyType guide0, guide1;  
    Node child0, child1, child2;  
}  
class LeafNode extends Node {  
    KeyType key;  
    ValueType value;  
}
```

Search(x):

$p \leftarrow$ root of tree, $h \leftarrow$ height of tree

repeat h times

// p points to an internal node

 if $x \leq p.\text{guide0}$ then $p \leftarrow p.\text{child0}$

 else if $p.\text{child2} = \text{null}$ or $x \leq p.\text{guide1}$ then $p \leftarrow p.\text{child1}$

 else $p \leftarrow p.\text{child2}$


// p now points to a leaf node

if $x = p.\text{key}$ then
 return $p.\text{value}$

else
 return null

A generalization: *B*-trees

- allow between B and $2B$ guides in each internal node
- branching factor is at least $B + 1$
- height of the tree is at most $\log_{B+1}(n)$
- example: $B = 2^{10}$ and $n = 2^{30}$, then height is just 3, instead of 30
- useful for high-latency memory (like hard drives)



- many file systems use *B*-trees to organize their metadata

Augmenting 2-3 trees

Idea: augment nodes with additional information to support new types of queries


Example: store # of keys in subtree at each internal node

Queries:


- What is the k th smallest key?
- How many keys are $\leq x$?

Fast range operations

- Suppose values associated with keys are numbers
- Operation $AddRange(x, y, \Delta)$ adds the same value Δ to the values associated with keys in the range $[x, y]$
- We can do this in time $O(\log n)$ using a “lazy” update technique
 - Store a value field at every node: internal nodes and leaves
 - “effective” value associated with a key is the sum all value fields on path from root to its leaf




- To perform *AddRange*(x, y, Δ):
 - trace paths e, f to x, y
 - add Δ to x, y , and to all roots of “internal” subtrees




- Insert and Delete operations also need to be slightly adjusted

Detailed Example:


(a) $\text{Search}(e) = 0+0+0+0+4=4$



(b) $\text{AddRange}(b, n, 1)$



(c) $\text{Search}(e) = 0+0+1+0+4=5$




2-3 Trees: Join and Split

$Join(T_1, T_2)$ joins two 2-3 trees in time $O(\log n)$

Assume $\max(T_1) < \min(T_2)$


Assume T_i has height h_i for $i = 1, 2$

Case 1: $h_1 = h_2$



Time: $O(1)$

Case 2: $h_1 < h_2$




- Attach v as the left-most child of p
- If p now has 4 children, we split p , and proceed up the tree as in *Insert*

Time: $O(h_2 - h_1 + 1) = O(\log n)$

Case 3: $h_1 > h_2$ — similar

$Split(T, x) \Rightarrow (T_1 [\leq x], T_2 [> x])$



We want to merge 2-3 trees X_1, \dots, X_k of heights h_1, \dots, h_k :

$$Y_1 := X_1$$

$$Y_2 := \text{Join}(Y_1, X_2)$$

$$Y_3 := \text{Join}(Y_2, X_3)$$

$$\vdots$$

$$Y_k := \text{Join}(Y_{k-1}, X_k)$$








Assumption:

- $h_i \leq h_{i+1}$ for $i = 1, \dots, k-1$,
- at most 2 trees of any given height — except the first 3 may be of the same height

Claim: Y_i has height h_i or $h_i + 1$ for $i = 2, \dots, k$

- *Proof:* See 2-3 Tree Handout

Example:



Claim \implies Time needed to compute
 $Y_{i+1} = \text{Join}(Y_i, X_{i+1})$ is $O(h_{i+1} - h_i + 1)$

\therefore the total cost is $O(t)$, where

$$\begin{aligned} t &\leq (h_2 - h_1 + 1) + \\ &\quad (h_3 - h_2 + 1) + \\ &\quad (h_4 - h_3 + 1) + \\ &\quad \vdots \\ &\quad (h_{k-1} - h_{k-2} + 1) + \\ &\quad (h_k - h_{k-1} + 1) \\ &= h_k - h_1 + k - 1 = O(h) \end{aligned}$$

and where h is the height of the original tree

Conclusion: total time for *Split* is $O(\log n)$